

DISCELL

Author One¹

Author Two²

¹Affiliation One ²Affiliation Two

Abstract

[TODO: abstract]

1 Introduction

Imaging-based spatial transcriptomics measures thousands of genes in every cell of a tissue section, in place (Janesick et al., 2023), so a cell’s expression can be read together with the cells around it. The counts assigned to a cell, however, come from three sources. The first is the cell’s *intrinsic state*: its type, its cell-cycle phase and whatever else it would carry anywhere in the tissue. The second is its *response* to its niche: programmes that change with who its neighbours are and what tissue surrounds it. The third is *spill-over*: transcripts of neighbouring cells assigned to it, because segmentation of densely packed tissue is imperfect. On this platform a typical cell carries about a tenth of its counts from neighbours (Yang et al., 2025).

The response and spill-over are confounded at first order. A cell that receives its neighbours’ transcripts acquires their markers; a cell that responds to its neighbours shifts its programmes in a way that their types predict. Both grow with the share of a neighbouring type (Bilous et al., 2026), and spill-over alone produces apparent cell-cell interaction signals (Mitchel et al., 2026). Where the responding genes are also markers of the neighbours, as for many ligands and receptors, the two leave the same signature. They differ in form (one adds a neighbour’s profile, the other shifts the cell’s own programmes), but from the counts alone the split between them is set by the modelling assumptions, and a flexible decoder can mimic either. Every claim that a spatial effect on expression is real therefore rests, usually silently, on how much of the neighbour-predictable signal was attributed to spill-over.

Related work. Classical models split a gene’s variance into spatial, intrinsic and environmental parts (Svensson et al., 2018; Arnol et al., 2019) or regress expression on niche covariates (Cable et al., 2022; Tanevski et al., 2022), and treat contamination separately, before analysis (Young and Behjati, 2020; Fleming et al., 2023; Petukhov et al., 2022; Mitchel et al., 2026). Among generative models of the cell, each models only one side of the confound (Table 1). On the spill-over side, resolVI (Ergen and Yosef, 2026) writes a cell’s counts as a mixture of its own decoded expression, its neighbours’ expression decoded by the same network, and a background, with the mixing weights estimated per cell. It has no niche term in the cell’s own expression, and it separates the components by expecting true expression to be low-dimensional. A niche response is low-dimensional too, and the part of it that moves a cell toward its neigh-

Table 1: The closest generative models, on the properties this paper turns on. *Spill-over*: misassigned transcripts are part of the generative model. *Invariance*: the regulariser that keeps the intrinsic latent free of the niche. *Intervention*: expression generated under an explicit change of the neighbourhood. ✓ yes, × no. ^aSIMVI instead estimates a spatial treatment effect.

Model	Niche	Spill-over	Invariance	Intervention
resolVI	×	estimated per cell	×	×
SIMVI	✓	×	independence, marginal adversarial, given type	×
MintFlow	✓	×	adversarial, marginal adversarial, given type	delete, replace cells
Cellina	✓	×	adversarial, marginal adversarial, given type	swap neighbours
DISCELL	✓	one fraction, swept	adversarial, given type	relocate a type

bours’ profile acts like spill-over, so that expectation alone does not separate the two. Our spill-over term is resolVI’s neighbour mixture with the weights fixed. On the niche side, NCEM (Fischer et al., 2023) regresses a cell’s expression on its neighbourhood given cell-type labels, including a conditional-VAE variant with a per-cell latent. SIMVI (Dong et al., 2025) splits a cell into an intrinsic latent and a spatially induced one, whose prior is a regression on the neighbours’ intrinsic latents, the closest analogue of our response prior, with an independence penalty applied to the intrinsic latent. MintFlow (Akbarnejad et al., 2025) conditions the intrinsic latent on cell type and keeps neighbour-type composition out of it with type-specific adversarial critics, the type-conditional invariance we also use. Cellina (Moeed et al., 2026) encodes an extrinsic latent from the neighbours’ expression, removes spatial-domain labels from the intrinsic latent adversarially, and predicts expression under a swapped neighbourhood. Celcomen (Megas et al., 2026) and NicheCompass (Birk et al., 2025) model inter-cellular gene programmes without a per-cell split. None of these models spill-over, and MintFlow and Cellina name segmentation errors as a limitation. SIMVI and MintFlow state identifiability results for their latents, but no method states what the counts cannot identify when a response and spill-over compete for them.

Our approach. DISCELL models both sides. A cell’s expected composition is a mixture of its own expression, built from an intrinsic state $\mathbf{z}_i$ and a response $\mathbf{w}_i$ whose prior depends on a descriptor of

the niche, and a fraction κ of its neighbours’ expression (Section 2.1). The counts bound κ from above but never from below (Proposition 1), so we do not estimate it. We refit the model over a grid of κ and report, for each finding, its *breakdown point* κ^* : the smallest spill-over fraction that explains the finding away (Definition 1). This is a sensitivity analysis in the tradition of tipping points and partial identification (Rosenbaum, 2002; Manski, 2003), with κ as the sensitivity parameter. Neither ingredient is new alone: the spill-over term is resolVI’s mixture with fixed weights, and adversarial invariance given type is MintFlow’s. What is new is putting a niche response and spill-over into one model, which creates the confound, and the sweep, which handles it.

Contributions.

- **A generative model in which a niche response and spill-over compete for the same counts** (Section 2.1).
- **Conditional invariance with a held-out probe, as an evaluation protocol.** A type-conditional adversary keeps neighbour composition and an image-derived niche descriptor out of $\mathbf{z}$; a probe trained after the fit measures the residual niche signal against a within-type permutation floor, for every method on the same cells (Section 2.3).
- **A spill-over sweep with breakdown points**, transferable to any model with a spill-over term, and with teeth: **[PENDING RUNS: breakdown-gaps queue: it breaks three findings, while the others hold to $\kappa = 0.4$, four times the typical misassignment]** (Section 2.4).
- **An evaluation on four sections** of ovarian cancer, lung cancer and paediatric lung, plus a serial section held out whole, against resolVI, SIMVI, MintFlow and Cellina, with recovery checks on simulated sections (Section 3). Code: **[URL to be added]**.

2 Model

DISCELL writes a cell’s counts as the sum of its intrinsic state, its response to the niche and spill-over from its neighbours; the first two are latent, the third is a fixed mixture whose size κ is swept (Figure 1).

2.1 Generative model

For each of N segmented cells i on one section we observe a count vector $\mathbf{x}_i$ over G genes with total ℓ_i ,

a lineage label $t_i \in \{1, \dots, K\}$, its position and the morphology image around it. Two cells are neighbours if their Voronoi cells share a face (Delaunay adjacency), with edges longer than $40 \mu\text{m}$ pruned; $\mathcal{N}_i$ is the neighbour set of i . The model is

$$\mathbf{z}_i \sim \mathcal{N}(\mathbf{0}, \mathbf{I}), \quad (1)$$

$$\mathbf{w}_i \mid \mathbf{c}_i, t_i \sim \mathcal{N}(m_\psi(\mathbf{c}_i, t_i), \sigma_w^2 \mathbf{I}), \quad (2)$$

$$\boldsymbol{\rho}_i = \text{softmax} \left(\underbrace{a(\mathbf{z}_i)}_{\text{intrinsic}} + \underbrace{\mathbf{B}\mathbf{w}_i}_{\text{response}} \right), \quad (3)$$

$$\bar{\boldsymbol{\rho}}_i = \sum_{j \in \mathcal{N}_i} \beta_{ij} \boldsymbol{\rho}_j, \quad (4)$$

$$\mathbf{p}_i = (1 - \kappa) \boldsymbol{\rho}_i + \underbrace{\kappa \bar{\boldsymbol{\rho}}_i}_{\text{spill-over}}, \quad (5)$$

$$\mathbf{x}_i \mid \ell_i \sim \text{Multinomial}(\ell_i, \mathbf{p}_i). \quad (6)$$

Here $\boldsymbol{\rho}_i$ is the cell’s own, *clean* composition over genes. Its intrinsic part decodes $\mathbf{z}_i \in \mathbb{R}^{d_z}$; its response part is linear in $\mathbf{w}_i \in \mathbb{R}^{d_w}$, so that column k of the loadings $\mathbf{B}$ is a gene programme and w_{ik} is how far the cell has moved along it. The decoder has no type input: the only route from identity to expression runs through $\mathbf{z}_i$, which therefore carries the type together with what the label discretises away, such as cell-cycle phase. The prior mean $m_\psi(\mathbf{c}_i, t_i)$ is the response expected of a cell of type t_i in niche $\mathbf{c}_i$; the prior scale is fixed at $\sigma_w = 1$, which pins the scale of $\mathbf{w}$ against that of $\mathbf{B}$.

The niche descriptor. $\mathbf{c}_i$ is a deterministic function of the data, shared by co-located cells and containing no latent:

$$\mathbf{c}_i = \text{GATv2}(\text{embed}(t_i); \{t_j\}_{j \in \mathcal{N}_i}) \oplus \Phi_i \oplus \mathbf{1}[\mathcal{N}_i = \emptyset]. \quad (7)$$

The first part is one graph-attention layer (Brody et al., 2022) over the neighbours’ types, queried with the cell’s own type, which amounts to a learned, type-specific reweighting of the neighbour-type composition $\mathbf{y}_i$. The second, Φ_i , is the embedding by a frozen pretrained image model (Shaban et al., 2025) of a $128 \mu\text{m}$ morphology crop around the cell, in which a disc of radius $25 \mu\text{m}$ around the cell is set to zero, so that it describes the surrounding tissue and not the cell. No distances enter, so that the attention cannot rediscover the spill-over kernel and merge the two mechanisms the model separates (Supplement S2.2).

Spill-over. The influx $\bar{\boldsymbol{\rho}}_i$ is the neighbours’ clean compositions averaged with a fixed kernel, $\beta_{ij} \propto f_{ij} \exp(-d_{ij}/\tau)$, normalised over $\mathcal{N}_i$, where d_{ij} is the distance between centroids, f_{ij} the length of the shared Voronoi face and $\tau = 20 \mu\text{m}$. The kernel




Figure 1: DISCELL in four panels. **a**, Spill-over: a cell (blue) receives transcripts from its neighbours in the pruned Delaunay graph, in proportion to the kernel weights β_{ij} (arrow widths); the niche descriptor $\mathbf{c}_i$ reads the neighbours’ types and the surrounding tissue image, with the cell itself masked out. **b**, The generative model (Equations (1) to (6)): the intrinsic state $\mathbf{z}_i$ and the response $\mathbf{w}_i$, whose prior is set by the niche and the type, decode to the cell’s own composition $\boldsymbol{\rho}_i$, which is mixed with the neighbours’ compositions at the fixed fraction κ . An adversary keeps the niche out of $\mathbf{z}_i$ given the type. **c**, Why κ is swept (Proposition 1): on the simplex of three genes, each assumed κ implies a clean composition on a line running from the observed composition $\mathbf{p}_i$ away from the neighbours’ influx $\bar{\boldsymbol{\rho}}_i$; every κ fits until a gene’s share would turn negative at $\bar{\kappa}_i$, so the data bound κ from above only. Schematic. **d**, The breakdown point (Definition 1): a finding is refitted at every κ of the grid, and κ^* is the smallest κ at which its interval reaches its null value. Schematic.

is computed once and not learned, unlike the per-cell neighbour weights of resolVI (Ergen and Yosef, 2025). The spill-over it represents has a specific form: one hop, one fraction for the whole section, and no gene dependence; an isolated cell receives none (Supplement S1.1).

2.2 Inference and objective

Both latents have amortised diagonal-Gaussian posteriors. The encoder of $q(\mathbf{z}_i \mid \mathbf{x}_i, t_i)$ reads the cell’s library-normalised log counts, its log total and its type, and never the niche. The posterior $q(\mathbf{w}_i \mid \mathbf{c}_i, t_i, \mathbf{z}_i, \mathbf{x}_i)$ reads the niche and also the cell, since $\mathbf{z}_i$ and $\mathbf{x}_i$ together carry evidence about $\mathbf{w}_i$ that the niche alone does not. Writing $r_i(\mathbf{w}) = \ell_i^{-1} \sum_g x_{ig} \log p_{ig}(\mathbf{z}_i, \mathbf{w})$ for the reconstruction per count, the training objective, maximised, is

$$J = \frac{1}{|\mathcal{S}|} \sum_{i \in \mathcal{S}} \left[r_i(\mathbf{w}_i) + \omega r_i(\tilde{\mathbf{w}}_i) - \alpha_a \text{Adv}_i - (1 + \omega) \alpha_z \text{KL}(q(\mathbf{z}_i) \parallel p(\mathbf{z}_i)) - \alpha_w \text{KL}(q(\mathbf{w}_i) \parallel p(\mathbf{w}_i \mid \mathbf{c}_i, t_i)) \right], \quad (8)$$

over a batch $\mathcal{S}$ of cells. With $\omega = 0$ and $\alpha_a = 0$ this is an evidence lower bound, scaled by depth so that the weights transfer between sections of different

depth. The second term is the *intrinsic path*: the same decoder evaluated with $\tilde{\mathbf{w}}_i$ drawn from the prior instead of the posterior, a second bound on the same evidence (Slavutsky et al., 2026) that asks $\mathbf{z}_i$ to explain the cell given only the response *expected* in its niche. It keeps the type out of $\mathbf{w}$: refitted with $\omega = 0$, the response’s agreement with type rises on all four sections (Supplement S8.2). The last term is the adversary of Section 2.3. The neighbours’ $\boldsymbol{\rho}_j$ enter $\bar{\boldsymbol{\rho}}_i$ under a stop-gradient, as in resolVI (Ergen and Yosef, 2025), and batches are contiguous spatial tiles with a two-hop halo (Supplement S3.3). We use $d_z = 20$, $d_w = 6$, $\omega = 1$, $\alpha_z = \frac{1}{2}/\bar{\ell}$ with $\bar{\ell}$ the section’s median total, and $\alpha_w = 0.1$; the remaining settings are in Supplement S3.1.

2.3 Conditional invariance

The architecture keeps the niche out of the encoder of $\mathbf{z}$, but not out of the counts it reads. The target is that, within a type, the intrinsic state says nothing about the niche:

$$I(\mathbf{z}_i; [\mathbf{y}_i, \Phi_i] \mid t_i) = 0. \quad (9)$$

It is conditional on the type because cell types cluster in space: “I am a T cell” predicts “my neighbours are T cells”, and an intrinsic state blind to the niche unconditionally would have to discard its type (Zhao and Gordon, 2019). It covers the image too, so that

a niche defined by morphology alone cannot hide in $\mathbf{z}$. Two heads with parameters ξ predict, from the posterior mean $\mu_{z,i}$ and the type, the neighbour-type composition $\mathbf{y}_i$ and a soft membership e_i^Φ of Φ_i in K image niches; the encoder is penalised by their skill beyond knowing the type alone,

$$\begin{aligned} \text{Adv}_i = \lambda_y [& \text{CE}(\mathbf{y}_i, \bar{\mathbf{y}}(t_i)) - \text{CE}(\mathbf{y}_i, \hat{\mathbf{y}}_\xi(\mu_{z,i}, t_i))] \\ & + \text{CE}(e_i^\Phi, \bar{\Phi}(t_i)) - \text{CE}(e_i^\Phi, \hat{e}_\xi^\Phi(\mu_{z,i}, t_i)), \end{aligned} \quad (10)$$

where $\bar{\mathbf{y}}(t)$ and $\bar{\Phi}(t)$ are the mean targets of type t and $\lambda_y = 3$ weights the composition, the channel of the spill-over confound. The heads maximise Adv_i with their own optimiser, the encoder minimises it, and it is zero at the optimum. Type-conditional adversarial invariance of an intrinsic latent against neighbour-type composition was introduced by MintFlow’s per-type critics (Akbarnejad et al., 2025); what we add is the image target and an independent test of the result.

The held-out probe. A training adversary that is defeated proves nothing (Elazar and Goldberg, 2018), so invariance is graded by a *fresh* predictor of each target block b (the composition, and the leading principal components of Φ) from $(\mu_{z,i}, t_i)$, trained on spatially separate training tiles and scored on held-out tiles. Its held-out gain over the type mean $\bar{v}_k(t)$, per component k of the block,

$$G_b = \frac{1}{|b|} \sum_{k \in b} \frac{1}{2} \log \frac{\mathbb{E} (v_k - \bar{v}_k(t))^2}{\mathbb{E} (v_k - \hat{v}_k(\mu_z, t))^2}, \quad (11)$$

estimates the predictive information from $\mathbf{z}$ to the block given the type (Xu et al., 2020). Its value with μ_z permuted among cells of the same type is the chance floor. The **residual niche signal** is the excess of G_b over that floor, reported as $1 - e^{-2 \text{ excess}}$: the share of the block’s within-type variance that the probe explains beyond chance. We fit a linear (ridge) and a nonlinear (small perceptron) probe, each against its own floor, and compare with a fit without the adversary ($\alpha_a = 0$). A probe at its floor bounds the dependence it can detect; it does not certify independence. Checkpoints are selected on held-out reconstruction while the agreement of $\mathbf{z}$ with type (normalised mutual information of a k -means clustering of μ_z with t) stays within 90% of its running maximum (Supplement S3.1).

2.4 What is identified

Spill-over and response both make a cell resemble its neighbours, and a flexible model can fit either.

The counts constrain the spill-over fraction in one direction only. Removing more and more of the neighbours’ profile from a cell eventually drives some gene’s share below zero, which no clean profile can have: this bounds κ from above. Every smaller κ yields a valid clean profile that reproduces the data exactly, so there is no lower bound (Figure 1c).

Proposition 1. *Hold the neighbours’ clean compositions, and hence $\bar{\rho}_i$, fixed, and let $\mathbf{p}_i = (1 - \kappa)\rho_i + \kappa\bar{\rho}_i$ with ρ_i in the simplex. A fraction $\kappa' \in [0, 1]$ reproduces $\mathbf{p}_i$ with a clean composition in the simplex if and only if $\kappa' \leq \bar{\kappa}_i = \min_{g: \bar{\rho}_{ig} > 0} p_{ig}/\bar{\rho}_{ig}$. The true κ satisfies $\kappa \leq \bar{\kappa}_i$, with equality if and only if $\rho_{ig} = 0$ for some gene g with $\bar{\rho}_{ig} > 0$.*

The proof is in Supplement S2.6. Even with the neighbours’ clean profiles known, then, the counts bound κ only from above, and pin it only through a gene that the cell cannot express itself. With finite counts the bound is soft; we do not compute it on our sections. In simulation, recovery of planted states and programmes barely changes with the assumed κ (Section 3.2), which is the same non-identifiability seen in practice. Given κ , the response is identified at most up to a rotation of its programmes and a per-type offset, so we report only rotation-invariant summaries and within-type contrasts (Supplement S1.5).

The sweep. We therefore do not estimate κ . We fit the model independently at each $\kappa \in \{0, 0.05, 0.1, 0.2, 0.3, 0.4\}$, with three seeds per value and every other weight held fixed, and check the held-out probe at each value, so that a change in a readout is not a change in the invariance. This is a sensitivity analysis with κ as the sensitivity parameter, in the tradition of tipping-point and E-value analyses of unmeasured confounding (Rosenbaum, 2002; VanderWeele and Ding, 2017): the fractions compatible with the data form an interval from zero (Manski, 2003), and the question asked of each finding is how far into it the finding survives.

Definition 1 (Breakdown point). For a signed contrast r , let $I_\kappa(r)$ be the interval of its mean over seeds at spill-over fraction κ . The contrast holds at κ if $I_\kappa(r)$ excludes zero and every seed keeps the sign the contrast has at $\kappa = 0$. Its *breakdown point* $\kappa^*(r)$ is the smallest grid value at which it does not hold. If it does not hold at $\kappa = 0$, r is no finding; if it holds at every grid value, $\kappa^*(r)$ is reported as exceeding the top of the grid.

The interval comes from a spatial block bootstrap over tiles, Bonferroni-adjusted over the contrasts claimed on a section. A finding with $\kappa^* > 0.4$ cannot

be explained by spill-over of the modelled form up to four times the typical misassignment of about a tenth of a cell’s transcripts (Yang et al., 2025); spill-over of another form (longer-ranged, gene-specific, depth-dependent) is not ruled out. Readouts shrink with κ by construction, so only sign and interval are read. A contrast without a sign at $\kappa = 0$, because it is zero there by construction or measures decoding rather than spill-over removal, is shown as a trajectory instead (Supplement S8.1). Nothing in the protocol is specific to DISCELL: any model with a spill-over term at a fixed fraction can attach a breakdown point to any spatial readout.

3 Experiments

3.1 Setup

[TODO: setup]

3.2 Simulation

[TODO: simulation]

3.3 What the Intrinsic State and the Response Carry

[TODO: intrinsic state and response]

3.4 What the Response Encodes

[TODO: programme atlas]

3.5 Comparison with Other Methods

Every method is scored by the same read-outs on its intrinsic latent, on our held-out cells (Section 3.1). Figure 2 sets two of them against each other: the residual niche signal (Section 2.3) and the within-type cell-cycle state the latent keeps. We use cycle state as the second axis because it is the one read of retained state that needs no label and varies within a cell type; it was fixed for all methods at once, after earlier baseline reads but before the final scoring. Agreement with the type label (NMI) is not a substitute: it rewards label fidelity, which MintFlow and Cellina are trained towards, and on it both score above DISCELL (Table S11).

No method both carries less residual niche signal than DISCELL and keeps more cycle state, on any section. Under a linear probe, DISCELL’s intrinsic state carries the least neighbour composition and the least image information of all methods on every section (on ovarian FFPE, 1.0–1.1% of the within-type

variance of composition, against 4.0% for MintFlow, the next lowest). Under the nonlinear probe of Figure 2, only MintFlow carries less wherever it ran, and its intrinsic latent keeps almost no within-type state (cycle $R^2 \leq 0.06$); on the TMA core, Cellina’s residual is at DISCELL’s lowest seed, with less cycle state. The methods that keep more cycle state carry more residual signal (Cellina on the two FFPE sections, SIMVI on the TMA sections), and on the fresh-frozen section DISCELL is best on both axes. DISCELL’s latent is also the one least predicted by the niche: the within-type R^2 of a linear map from DISCELL’s niche descriptor $\mathbf{c}_i$ to each method’s intrinsic latent (mirror R^2) is lowest for DISCELL on four of the five sections, MintFlow being lower on the serial one. [PENDING RUNS: MintFlow refits: the MintFlow comparisons in this paragraph are re-scored]

resolVI keeps a single latent for the cell, so it is a reference for a model that does not separate the niche, not a competitor on this axis; its residual is two to five times DISCELL’s on the three whole-slide sections. The result is not reproduced by adding an adversary to another model. Given our niche label as the domain its adversary removes, Cellina’s best case on this probe, its residual does not consistently fall, and it loses agreement with type and cycle state on every section; on its own neighbour graph it keeps less cycle state than DISCELL at a similar or higher residual (Supplement S4.3). The separation comes from the architecture, not from the adversary alone.

Held-out reconstruction is within 0.015 nats per count of the best method on every section and best on the two whole-slide FFPE sections, although the comparison methods were trained on these held-out cells (Table S11). One DISCELL fit, its evaluations included, is up to 2.6 times faster than the fastest comparison method, Cellina, on three of the four sections and slower than it on the fresh-frozen section; it is over a hundred times faster than MintFlow wherever MintFlow could be run (Table S19). Because the result is a sweep of 18 fits (six spill-over fractions, three seeds; Section 2.4), the honest unit is the sweep: at the operating point’s run time it takes at most 2.8 hours on each section where MintFlow ran, less than one MintFlow fit there (9 to 16 hours), and on the TMA core less than one SIMVI fit; on the fresh-frozen section it takes longer than one resolVI fit. [PENDING RUNS: MintFlow refits: MintFlow’s reconstruction and run times] The full battery, probe and context-side tables are in the supplement (Tables S6, S11 and S12).




Figure 2: Residual niche signal in the intrinsic latent against the cell-cycle state it keeps, per section. *Horizontal*: share of the within-type variance of neighbour composition that the held-out nonlinear probe reads from each method’s intrinsic latent beyond the permutation floor, in %. *Vertical*: R^2 of the cell-cycle scores from the intrinsic latent on the most strongly cycling tenth of held-out cells. The desirable corner is the top left. DISCELL: mean over three seeds, bars the range; open circle: DISCELL without the adversary. resolVI keeps a single latent and is shown as a reference. Cellina is also shown with our niche label as the domain its adversary removes (its best case on this probe), and on the TMA sections on its own neighbour graph. SIMVI ran out of memory on a 24 GB GPU on the three larger sections, and MintFlow could not be fitted on the fresh-frozen section; both are shown where their whole-section fits exist. The values are in Tables S6 and S11. [PENDING RUNS: MintFlow refits: MintFlow’s points are redrawn from the new fits]

3.6 Relocation

Relocation asks how a cell type’s expression would change if its cells sat in another niche. DISCELL answers it by keeping each cell’s intrinsic state, moving the prior of its response and its spill-over influx to the target niche, and decoding. For each type and pair of niches (a panel), the predicted change in the type’s mean expression is scored against the observed difference between the two niches, as a share of the part of that difference that is reproducible between two random halves of the cells (the noise ceiling); only panels whose ceiling is high enough to be read count (Supplement S5.1).

Moved to another niche, a cell type’s predicted shift recovers 0.65 to 0.73 of the reproducible between-

niche difference (??). The response programmes of Section 3.4 carry most of it: predicted from the programmes alone, the shift recovers 0.52 to 0.64 of the ceiling, and from the spill-over alone 0.21 to 0.39 (the two parts are not additive; Table S17).

Cellina (Moeed et al., 2026) also predicts expression under a rewired neighbourhood, so we score its counterfactual by the same reads on the same panels (Table 2). The comparison is mixed. DISCELL’s moved cells are closer to the target cells’ own distribution on every section, with disjoint intervals on three, and DISCELL leads on all three reads on the fresh-frozen section. Cellina recovers more of the ceiling on the ovarian FFPE section, and on the TMA core on its single readable panel, and is narrowly

Table 2: Relocation, DISCELL against Cellina’s neighbour-rewiring counterfactual, on the same panels, niches and ceilings (DISCELL’s first seed, which supplies the panels; Cellina refitted on the training tiles). *Ceiling*: share of the reproducible between-niche difference recovered, over the panels whose ceiling can be read. *Single cell*: share of the distribution gap to the target niche’s own cells that the moved cells close. *Twin*: how much closer, in expression, each target cell lies to the moved cell nearest it in intrinsic state than to a random moved cell. D: DISCELL; C: Cellina. Higher is better; bold marks the better of the two. ^aA single readable panel. Intervals: Table S17.

Section	Ceiling		Single cell		Twin	
	D	C	D	C	D	C
Ovarian FFPE	0.650	0.731	0.789	0.681	0.398	0.361
Lung FFPE	0.743	0.703	0.826	0.757	0.432	0.435
Ovarian FF	0.677	0.566	0.823	0.602	0.437	0.330
TMA core	0.711	0.899^a	0.814	0.801	0.464	0.284

ahead on the twin read on lung. Most scored cells lie in the model’s training tiles; scored on held-out tiles only, DISCELL’s share of the ceiling moves in both directions across fits, so it shows no systematic in-sample advantage, while Cellina’s falls on both FFPE sections, on lung from 0.70 to 0.49, below DISCELL’s range (Table S15).

Would a plain regression do as well? Regressing each type’s expression shift directly on neighbour composition, from the counts as observed, recovers [PENDING RUNS: plain-regression reference: share of the ceiling and single-cell read on the same panels] of the ceiling on the same panels. [PENDING RUNS: keep one of the two planned sentences once the reference lands] *If below DISCELL*: modelling spill-over and keeping the intrinsic state apart also makes the relocation more accurate than the niche composition alone predicts. *If level with DISCELL*: the relocation’s accuracy then comes from the niche composition itself, and DISCELL’s advantage lies in the separation, an intrinsic state largely free of the niche and spill-over modelled rather than absorbed, not in relocation accuracy.

3.7 Robustness to Spill-over

Every section was refitted along the sweep of Section 2.4, and Table 3 gives each finding of Sections 3.3 to 3.6 its breakdown point: how much spill-over, assumed in the model, it takes to explain the finding away.

Table 3: Breakdown points κ^* : the smallest assumed spill-over fraction at which a finding stops holding, per section, on the grid $\kappa \in \{0, 0.05, 0.1, 0.2, 0.3, 0.4\}$ with three seeds each. >0.4 : the finding holds at every grid point. Bold: a break. LR: ligand–receptor. ^aThe interval of the seed mean contains 0; ^ba seed has the opposite sign. n/a: no tumour cells on the TMA sections. The serial section is read through the core’s fits. ?: [PENDING RUNS: breakdown gaps, landing tonight; findings per section become 7/7/7/6/6]. Trajectories and diagnostics: Supplement S8.1.

Finding	Ov. FFPE	Lung FFPE	Ov. FF	TMA core	TMA serial
<i>Biological claims</i>					
Relocation > spill-over part	0.4^b	0.3^b	>0.4	>0.4	?
Relocation > type mean	>0.4	>0.4	>0.4	>0.4	?
Nearest-twin advantage	>0.4	>0.4	>0.4	>0.4	?
LR-gene lean of w	>0.4	>0.4	0.2^a	>0.4	?
Tumour axis in w	>0.4	?	?	n/a	n/a
<i>Allocation checks</i>					
Cycle in z , not in w	>0.4	>0.4	>0.4	>0.4	>0.4
Niche information in w	>0.4	>0.4	>0.4	>0.4	?

The test has teeth. It breaks three findings: the response’s lean toward ligand–receptor genes on the fresh-frozen section ($\kappa^* = 0.2$), and the relocation’s margin over its spill-over part on the ovarian and lung FFPE sections ($\kappa^* = 0.4$ and 0.3 , where a seed changes sign). These are the findings that spill-over itself could most plausibly produce: a gene-specific spill-over would create the lean, and the margin measures what the prediction adds to spill-over. Every other finding holds across the whole grid on every section, up to four times the typical misassignment. [PENDING RUNS: breakdown gaps: the tumour axis on lung FFPE and ovarian FF and five findings on the serial section; recount the breaks] The table lists the allocation checks, that **z** keeps the cycle and **w** the niche, apart from the biological claims, because the split favours them by design; the biological claims are the stronger test. Two contrasts first in the graded set were moved to trajectories after reading, for the reasons of Section 2.4: one, between marker pairs, measures decoding, and the other, the relocation’s margin over its programme part, is zero by construction at $\kappa = 0$.

A positive control shows that the rule detects spill-over when it is there. In simulation, we plant a spatial contrast created only by spill-over at a known fraction, beside a genuine response contrast in the same world, and run the same sweep. The prediction, set before the run, is that the spill-over contrast

breaks at the first grid point at or above the planted fraction, while the response contrast holds across the grid. **[PENDING RUNS: planted spill-over control: planted fraction, observed κ^* , and the response contrast’s status; if κ^* does not track the planted fraction, report it as found and present the breakdown point as a lower-power check]**

Of the spill-over forms the sweep leaves open (Section 2.4), several were refitted on the primary section, three seeds each (Supplement S8.2): per-cell and per-gene spill-over rates move no read by more than 1.3 seed standard deviations, but a fixed false-positive floor lowers the relocation read from 0.65 to 0.50.

4 Discussion

[TODO: discussion]

References

- Akbarnejad, A., Steele, L., Jafree, D. J., Birk, S., Sallese, M. R., et al. (2025). Mapping and reprogramming human tissue microenvironments with MintFlow. *bioRxiv*.
- Arnol, D., Schapiro, D., Bodenmiller, B., Saez-Rodriguez, J., and Stegle, O. (2019). Modeling Cell-Cell Interactions from Spatial Molecular Data with Spatial Variance Component Analysis. *Cell Reports*, 29(1):202–211.e6.
- Bilous, M., Buszta, D., Bac, J., Kang, S., Dong, Y., Tissot, S., Andre, S., Alexandre Gaveta, M., Voize, C., Peters, S., Homicsko, K., and Gottardo, R. (2026). Resolving sensitivity, specificity and signal contamination in Xenium spatial transcriptomics. *Nature Methods*, 23:1152–1162.
- Birk, S., Lotfollahi, M., et al. (2025). Quantitative characterization of cell niches in spatially resolved omics data. *Nature Genetics*. CHECK: author list, volume, identifier.
- Brody, S., Alon, U., and Yahav, E. (2022). How attentive are graph attention networks? In *International Conference on Learning Representations*.
- Cable, D. M., Murray, E., Shanmugam, V., Zhang, S., Zou, L. S., Diao, M., Chen, H., Macosko, E. Z., Irizarry, R. A., and Chen, F. (2022). Cell type-specific inference of differential expression in spatial transcriptomics. *Nature Methods*, 19(9):1076–1087. Preprint: doi:10.1101/2021.12.26.474183.
- Dong, M., Kluger, Y., et al. (2025). SIMVI disentangles intrinsic and spatial-induced cellular states in spatial omics data. *Nature Communications*. CHECK: author list, volume, identifier.
- Elazar, Y. and Goldberg, Y. (2018). Adversarial removal of demographic attributes from text data. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 11–21.
- Ergen, C. and Yosef, N. (2025). ResolVI – addressing noise and bias in spatial transcriptomics. *bioRxiv* preprint.
- Ergen, C. and Yosef, N. (2026). ResolVI: addressing noise and bias in spatial transcriptomics. *Nature Methods*.
- Fischer, D. S., Schaar, A. C., and Theis, F. J. (2023). Modeling intercellular communication in tissues using spatial graphs of cells. *Nature Biotechnology*, 41:332–336.
- Fleming, S. J., Chaffin, M. D., Arduini, A., Akkad, A.-D., Banks, E., Marioni, J. C., Philippakis, A. A., Ellinor, P. T., and Babadi, M. (2023). Unsupervised removal of systematic background noise from droplet-based single-cell experiments using CellBender. *Nature Methods*, 20(9):1323–1335. Preprint: doi:10.1101/791699.
- Janesick, A., Shelansky, R., Gottscho, A. D., Wagner, F., Williams, S. R., Rouault, M., Beliakoff, G., Morrison, C. A., Oliveira, M. F., Sicherman, J. T., et al. (2023). High resolution mapping of the tumor microenvironment using integrated single-cell, spatial and in situ analysis. *Nature Communications*, 14:8353. CHECK: full author list truncated.
- Manski, C. F. (2003). *Partial Identification of Probability Distributions*. Springer.
- Megas, S., Chen, D. G., Polanski, K., Asadolahzadeh, H., Eliasof, M., Schönlieb, C.-B., and Teichmann, S. A. (2026). Celcomen: spatial causal disentanglement for single-cell and tissue perturbation modeling. *Nature Communications*, 17:4126.
- Mitchel, J., Gao, T., Petukhov, V., Cole, E., and Kharchenko, P. V. (2026). Impact and correction of segmentation errors in spatial transcriptomics. *Nature Genetics*, 58:434–444.
- Moeed, A., Schrod, S., Rohbeck, M., Bonder, M. J., Lutsik, P., Stegle, O., and Dimitrov, D. (2026). Querying counterfactuals on tissue graphs with supervised disentanglement. *arXiv preprint* 2606.08493.

- Petukhov, V., Xu, R. J., Soldatov, R. A., Cadinu, P., Khodosevich, K., Moffitt, J. R., and Kharchenko, P. V. (2022). Cell segmentation in imaging-based spatial transcriptomics. *Nature Biotechnology*, 40:345–354.
- Rosenbaum, P. R. (2002). *Observational Studies*. Springer, 2nd edition.
- Shaban, M., Chang, Y., Qiu, H., Yeo, Y. Y., Song, A. H., Jaume, G., et al. (2025). A foundation model for spatial proteomics. arXiv preprint 2506.03373.
- Slavutsky, Y., Beker, O., Blei, D., and Dumitrescu, B. (2026). Variational learning of disentangled representations. In *International Conference on Machine Learning*. arXiv:2506.17182.
- Svensson, V., Teichmann, S. A., and Stegle, O. (2018). SpatialDE: identification of spatially variable genes. *Nature Methods*, 15:343–346.
- Tanevski, J., Ramirez Flores, R. O., Gabor, A., Schapiro, D., and Saez-Rodriguez, J. (2022). Explainable multiview framework for dissecting spatial relationships from highly multiplexed data. *Genome Biology*, 23:97.
- VanderWeele, T. J. and Ding, P. (2017). Sensitivity analysis in observational research: Introducing the E-value. *Annals of Internal Medicine*, 167(4):268–274.
- Xu, Y., Zhao, S., Song, J., Stewart, R., and Ermon, S. (2020). A theory of usable information under computational constraints. In *International Conference on Learning Representations*.
- Yang, Y., DePasquale, E., Adeleke, D., Xie, X., Lu, Z., Nelson, M. D., Xiao, G., Wang, X., and Wang, Y. (2025). MisTIC: Missegmented transcript inference correction for improved spatial transcriptomics analysis. bioRxiv preprint.
- Young, M. D. and Behjati, S. (2020). SoupX removes ambient RNA contamination from droplet-based single-cell RNA sequencing data. *GigaScience*, 9(12):giaa151.
- Zhao, H. and Gordon, G. J. (2019). Inherent tradeoffs in learning fair representations. In *Advances in Neural Information Processing Systems*, volume 32.